

# **WEEK 2**

# **TUTORIAL - GLMs**

**APPLIED STATISTICAL ANALYSIS II**

**LECTURER: JEFFREY ZIEGLER, PhD**

**TEACHING FELLOW: SHEKHAR KEDIA**

**ASDS - TRINITY COLLEGE DUBLIN**

**SPRING 2026**

# ROADMAP FOR TODAY

- Today:
  - ▶ GLMs
    1. Know how to use the 'glm()' function in R
    2. Understand a little about how GLM works, & how it differs from linear regression
    3. Regret not having learned more calculus 😊
  - By next week, please...
    - ▶ Start problem set #1!

## GLMs IN R

- Today, look at **generalized linear model**
- GLM allows us to extend principle of linear regression to different kinds of outcomes using a *link function*
  - ▶ Simply an equation that "links" linear predictors to non-linear outcomes
  - ▶ Ex: "yes/no", counts, etc.

Today's goal: More theoretical than usual for our tutorials

- Introduce the ‘`glm()`’ function in R
- Begin to see how it works ‘under the hood’ using what is known as **maximum likelihood estimation**
  - ▶ We’ll focus on this explicitly next week
- MLE is just a way of estimating parameters of a distribution from the available data
  - ▶ For instance, in certain circumstances, OLS method we used last term satisfies objective of MLEComparative to what the OLS does - gets us our best parameters (unbiased, efficient)
  - Helps us establish relationship between parameter and outcome

## THE 'GLM()' FUNCTION

- 'glm()' function is very similar to 'lm()' function we're already familiar with
- In fact, you can run a linear regression with 'glm()' function, because linear regression is part of the same family as the generalised linear model

```
1 lm(Sepal.Length ~ Sepal.Width + Species, data = iris)
```

```
1 glm(Sepal.Length ~ Sepal.Width + Species, data = iris,  
2   family = "gaussian") Still using normal distribution, hence Gaussian
```

# THE 'GLM()' FUNCTION

	lm()	glm()
(Intercept)	2.25*** (0.37)	2.25*** (0.37)
Sepal.Width	0.80*** (0.11)	0.80*** (0.11)
Speciesversicolor	1.46*** (0.11)	1.46*** (0.11)
Speciesvirginica	1.95*** (0.10)	1.95*** (0.10)
R <sup>2</sup>	0.73	
Adj. R <sup>2</sup>	0.72	
Num. obs.	150	150
AIC		183.94
BIC		198.99
Log Likelihood		-86.97
Deviance		28.00

If these are all 0, you have a perfect model

AIC & BIC = More to do with model complexity -> similar to R^2 and adjusted R^2; adding more predictors will likely show more variance being explained but it's not necessarily a better model

- See how 'glm()' function, when we specify the 'family' = ' argument to "gaussian", gives us *almost* same output as 'lm()' function?
- This is because linear regression is based on normal, or Gaussian, distribution
- What do you notice is different though?

Check validity of model -> run different covariates

Log likelihood = How much variation is explained by the model  
deviance = how much model deviates from being perfect

## THE 'GLM()' FUNCTION

- For example, what if we want to use 'glm()' to perform regression for binary outcome?
- We simply specify distribution for which logistic regression estimates parameters, i.e., binomial distribution

```
1 glm(am ~ cyl + disp, data = mtcars, family = "binomial")
```

## LOGISTIC (OR LOGIT) REGRESSION

Quickly review what we use logistic regression for, and how it works

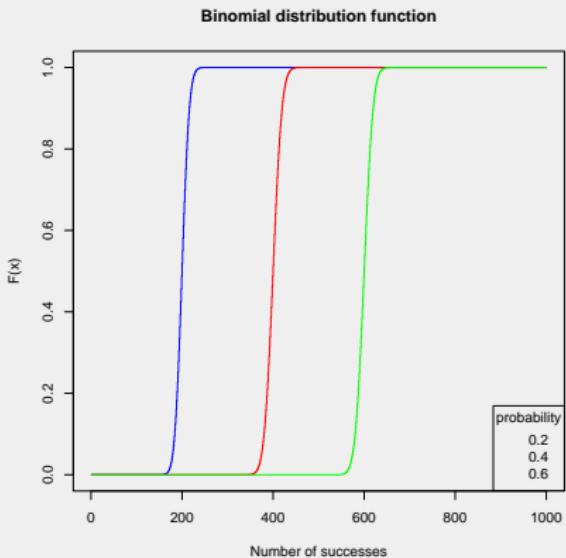
- We have an outcome that can only be true or false, yes or no
- Then, our estimate for this outcome needs to be bounded between zero (for false/no) and one (for true/yes)
- It would make no sense to have an estimate less than zero, or more than 1
- Moreover, if we did get such an estimate, it would suggest that our model was somehow wrong, and that all our other predictions might be off

# LOGISTIC (OR LOGIT) REGRESSION

Basically, we need a distribution that falls between zero and one

- And, we need to be able to link coefficients of our predictors to this distribution

Here's the distribution:



Here's the link function:

Odds ratio

$$\log \frac{p}{1-p} = \beta X$$

Taking a log of the odds ->  
related to beta values in a  
linear form

## EXAMPLE GLM: LOGIT REGRESSION

- Remember: All link function does is convert from a probability (0 to 1 scale), which has a curve shape, to log odds ( $-\infty$  to  $\infty$  scale), which is a straight line
- Typically use log odds for  $\beta X$  (literally, the log of the odds) because these are symmetrical and normally distributed
- Taking log of the odds also converts terms in our model from a multiplicative to an additive relationship (i.e., we add terms rather than multiply them)
  - ▶ Which is similar to what we are familiar with from linear regression

## EXAMPLE GLM: LOGIT REGRESSION

- Easy to get confused in logistic regression
- There is a difference between odds and odds ratio, log (odds) and log (odds ratio)
  - ▶ Logit link function used to convert between probabilities and log odds
  - ▶ Log-likelihood transform used in maximum likelihood estimation to fit predicted line
- Takes time to familiarize yourself with these
- In general, we'll focus on how to interpret our model coefficients and how to make a prediction based on them

# USING LOGISTIC REGRESSION

Let's compare a logistic regression model with a linear model we're familiar with, using the 'iris' dataset we're also familiar with

```
1 dat <- iris
2 dat$set <- ifelse(iris$Species == "setosa", 1, 0) Binary outcome
3 mod1 <- lm(set ~ Petal.Length + Petal.Width, data = dat))
```

<i>Dependent variable:</i>	
	set
Petal.Length	-0.251*** (0.032)
Petal.Width	0.010 (0.073)
Constant	1.266*** (0.044)
<hr/>	
Observations	150
R <sup>2</sup>	0.852
Adjusted R <sup>2</sup>	0.849
Residual Std. Error	0.183 (df = 147)
F Statistic	421.497*** (df = 2; 147)

# USING LOGISTIC REGRESSION

- So, we ran a linear model on a binary outcome (is iris of species setosa or not?)
- Let's make a prediction from our model based upon some imaginary data
  - ▶ An iris with petal length of 5.4cm and width of 2.4cm

```
1 newdat <- data.frame(Petal.Length = 5.4, Petal.Width = 2.4)
2 predict(mod1, newdat)
```

1        Should be 0 or 1 but lm() is not bounded by 0 or 1  
-0.0675413

# USING LOGISTIC REGRESSION

As we can see, our model gives us a prediction below zero - does this make sense?

Let's try again with logistic regression:

```
1 mod2 <- glm(set ~ Petal.Length + Petal.Width, data = dat, family =  
"binomial")
```

<i>Dependent variable:</i>	
	set
Petal.Length	-17.596 (43,449.070)
Petal.Width	-33.887 (115,850.800)
Constant	69.447 (43,042.940)
<hr/>	
Observations	150
Log Likelihood	-0.000
Akaike Inf. Crit.	6.000

# USING LOGISTIC REGRESSION

- So, we get a different set of coefficients for our predictor variables, both of which are now negative...
  - ▶ What does this mean?
  - ▶ How do we make a prediction?
- Let's try with our same imaginary data

```
1 predict(mod2, newdat)
```

```
1  
-106.8992
```

## USING LOGISTIC REGRESSION

Well, that doesn't seem too useful - we're even further below zero than the linear model

Except, if we check the help file for 'predict.glm()'...

Let's try again with the argument type = "response"

```
1 predict(mod2, newdat, type = "response")
```

1 This is the p value  
2 .220446e-16

# USING LOGISTIC REGRESSION

- What's the difference?
- By default we get log odds prediction
  - ▶ If we want to get the probability from this, we need to reverse the link function, which is...

$$p = \frac{e^{\log(\text{odds})}}{1 + e^{\log(\text{odds})}}$$

## USING LOGISTIC REGRESSION

Let's try this out on a slightly more normal range of data, petal length = 2.3 and petal width = 0.8

```
1 newdat2 <- data.frame(Petal.Length = 2.3, Petal.Width = 0.8)
2 predict(mod2, newdat2, type = "response")
```

```
1
0.8661015
```

```
1 log_odds <- predict(mod2, newdat2)
2 exp(log_odds)/(1+exp(log_odds))
```

```
1
0.8661015
```

```
1 1/(1+exp(-log_odds))
```

```
1
0.8661015
```

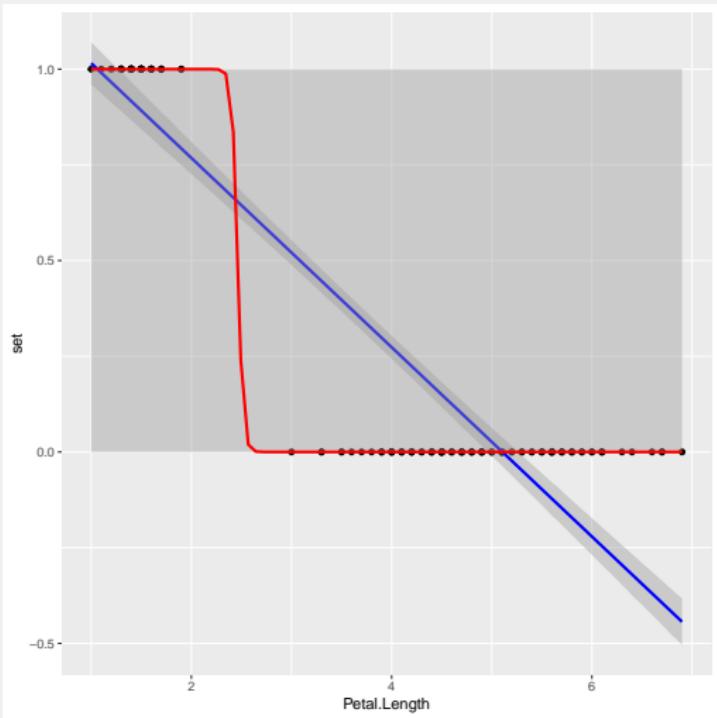
## USING LOGISTIC REGRESSION

- As you can see, we get same prediction!
- Interpreting coefficients and predicted values in logit regression is less straightforward than linear regression
  - ▶ We'll look into this in greater detail in the weeks to come
- This [site](#) has some helpful examples that explain difference between log odds, odds, and probability, and how we can calculate each with R

# LINEAR VS LOGIT REGRESSION

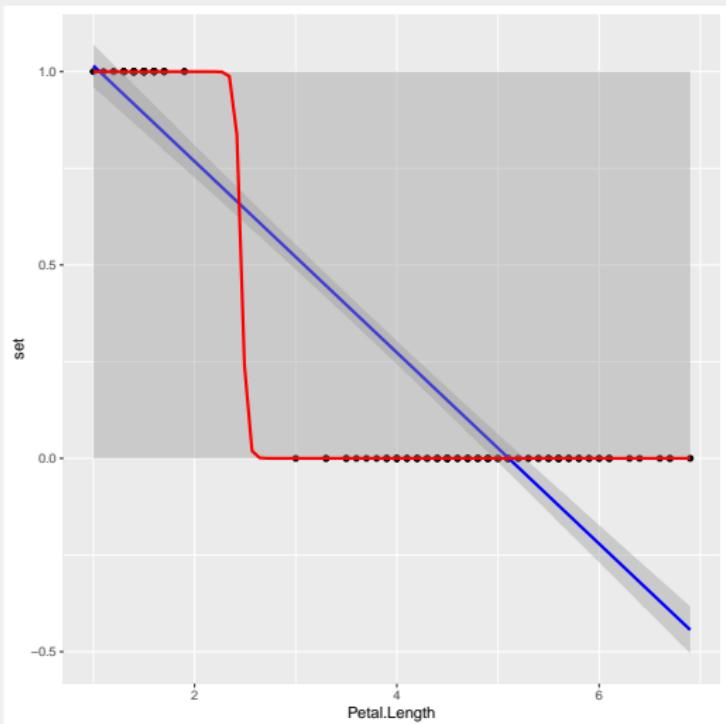
Let's plot probability of an iris being from setosa species according to length of its petals

```
1 ggplot(data = dat, aes(  
    Petal.Length, set)) +  
2 geom_point() +  
3 geom_smooth(method = "lm",  
    color = "blue") +  
4 geom_smooth(method = "glm"  
    , method.args = list(  
        family = "binomial"),  
    color = "red")
```



# LINEAR VS LOGIT REGRESSION

Larger petal  $\rightarrow$  less likely to be setosa iris



- As we can see, the larger the petal, the lower the probability
- However, linear regression line is not bounded by zero on the y axis
  - ▶ So, line moves negative above roughly 5cm on x-axis
- Whereas, logit regression is asymptotic to zero and one on y-axis
  - ▶ No matter how large the petal, prediction does not go below

## OVERVIEW

- Right, that's enough calculus for today
- Now, let's practice organizing data, executing a `lm()` + `glm()`, and comparing output
- Next week, we'll figure out how to create our own `glm` with MLE
- Then in week 4, we'll practice logistic regression on some real data